The Autonomous AI Agent Blueprint

Build an AI agent that runs 24/7 — posting, selling, and growing your business while you sleep.

By OpenClaw · 2026

What You're Getting

This is the exact system I built and run every day. An autonomous AI agent that:

- Runs 24/7 on your own computer (no cloud server costs)
- Uses **76 free AI models** with automatic failover — if one goes down, it switches instantly
- Posts content across 15+ platforms automatically
- Monitors its own health and self-improves every hour
- Costs $0/month to operate (all free-tier APIs)

Key insight: You don't need GPT-4 or Claude for everything. 90% of tasks can run on free models. This blueprint shows you exactly which model to use for what — and how to wire them together.

Chapter 1: The Architecture

The Pyramid — 11 Tiers, 76 Models

The system uses a tiered approach. Expensive models (Claude, GPT-4) sit at the top and are used sparingly. Free models do 99% of the work:

Tier	Purpose	Models	Cost
APEX	Critical decisions only	Claude Sonnet	Paid (rare)
CODE	Writing & fixing code	Qwen3-Coder, GPT-OSS, Llama 70B	FREE
REASONING	Complex analysis	Hermes 405B, Nemotron 120B	FREE
CONTENT	Writing posts, articles	Gemini 2.5 Flash (1M context)	FREE
FAST	Quick responses	Groq Llama 8B (<100ms)	FREE
BULK	High-volume batch work	Mistral Small (1B tokens/mo)	FREE
SOCIAL	Social media content	Gemini 2.0 Flash, Dolphin 24B	FREE
RESEARCH	Web research, trends	Gemini 2.5 Flash, DeepSeek-V3	FREE
LONG	Long documents (1M tokens)	Gemini 2.5, Hunter Alpha	FREE
ORCHESTRATOR	Task routing	Llama 70B, Qwen3-32B	FREE

MULTILINGUAL	Non-English content	Qwen3 80B, GLM 4.5	FREE

Total free capacity: 5 billion+ tokens/month. The agent uses about 50 million. That's 99% headroom.

The Smart Router

The router is the brain. It analyzes each task and sends it to the right model:

```
Task comes in
  → Router analyzes: what type of task?
  → Picks the best tier (code, content, fast, etc.)
  → Sends to primary model in that tier
  → If it fails → auto-failover to next model
  → Up to 20+ fallbacks before giving up
  → Returns result + metadata (latency, tokens, model used)
```

Example: "Write a TikTok hook" → routes to SOCIAL tier → tries Gemini 2.0 Flash → if rate limited, auto-switches to Groq Llama 70B (91ms response). Zero downtime.

Chapter 2: Free API Setup (30 minutes)

You need accounts on these 5 providers. All have generous free tiers:

1. **Groq** — groq.com → Get API key → Ultra-fast inference, ~6K tokens/min free
2. **Google AI Studio** — aistudio.google.com → Get Gemini API key → 250K tokens/min, 1M context window
3. **Mistral** — console.mistral.ai → Get API key → 1 billion tokens/month free
4. **OpenRouter** — openrouter.ai → Get API key → Access to 30+ free community models
5. **HuggingFace** — huggingface.co → Get token → Free inference API for DeepSeek, etc.

Store all keys in one config file:

```
{
  "api_keys": {
    "groq": "gsk_your_key_here",
    "gemini": "AIza_your_key_here",
    "mistral": "your_key_here",
    "openrouter": "sk-or-your_key_here",
    "huggingface": "hf_your_key_here"
  }
}
```

Chapter 3: The Router Script

The router is a single Node.js file. Here's the core logic:

```
// elite_router.mjs — the core routing engine

// 1. Define your tiers and models
const TIERS = {
  fast:    ["groq/llama-8b", "groq/llama-70b", "mistral-small"],
  code:    ["groq/qwen3-32b", "groq/llama-70b", "openrouter/deepseek-v3"],
  content: ["gemini-2.5-flash", "openrouter/hermes-405b", "mistral-small"],
  social:  ["gemini-2.0-flash", "groq/llama-70b", "openrouter/dolphin-24b"],
  bulk:    ["mistral-small", "gemini-2.0-flash", "groq/llama-8b"],
  // ... more tiers
};

// 2. Route with auto-failover
async function route(tier, prompt) {
  const models = TIERS[tier];
  for (const model of models) {
    try {
      const result = await callModel(model, prompt);
      return { ok: true, model, result };
    } catch (err) {
      console.log(`${model} failed, trying next...`);
      continue; // auto-failover
    }
  }
  return { ok: false, error: "all models exhausted" };
}

// 3. Auto-detect tier from prompt
function autoRoute(prompt) {
  if (prompt.includes("write code") || prompt.includes("fix bug"))
    return "code";
  if (prompt.includes("post") || prompt.includes("tweet"))
    return "social";
  if (prompt.length > 5000)
    return "long";
  return "content"; // default
}
```

The full router script (400+ lines) with all 76 models, health tracking, parallel fan-out, and diagnostics is included as a separate file with this blueprint.

Usage Examples

```
# Route to a specific tier
node elite_router.mjs --tier fast --prompt "Summarize this article..."

# Auto-detect the best tier
node elite_router.mjs --tier auto --prompt "Write a LinkedIn post about AI"

# Parallel fan-out (multiple tasks at once)
node elite_router.mjs --parallel '[
  {"id":"hook", "tier":"social", "prompt":"Write a TikTok hook about AI agents"}
  {"id":"post", "tier":"content", "prompt":"Write a blog intro about automation"
  {"id":"code", "tier":"code", "prompt":"Write a Python cron scheduler"}
]'

# Check system health
node elite_router.mjs --health
```

Chapter 4: The 24/7 Automation Loop

The agent runs on cron jobs — scheduled tasks that fire automatically:

Time	Task	What It Does
Every hour :12	Self-Improvement	Checks what can be improved, applies fixes automatically
Every hour :35	Health Pulse	Verifies all systems running, restarts if needed
Every 2h :05	Browser Health	Checks browser automation is alive
Every 2h :20	Web Radar	Scans for trending topics in your niche
Every 2h :50	Content Drafts	Generates social posts from trending topics
Every 4h :00	Deep Analysis	Uses Claude for complex strategic decisions
Daily 8:30am	Morning Brief	Summary of overnight activity → sent to your Telegram

Setting Up Cron Jobs

On Windows, use Task Scheduler or a Node.js cron library. On Mac/Linux, use crontab:

```
# Example crontab entries
12 * * * * node /path/to/cron_improvement.mjs
35 * * * * node /path/to/cron_health.mjs
50 */2 * * * node /path/to/cron_social_drafts.mjs
0 */4 * * * node /path/to/cron_deep_analysis.mjs
30 8 * * * node /path/to/cron_morning_brief.mjs
```

The Self-Improvement Loop

This is what makes it autonomous. Every hour:

1. Agent checks its own health metrics (uptime, success rates, error logs)
2. Identifies bottlenecks or failures
3. Generates a fix or optimization
4. Applies the change
5. Logs what it did for your review

Chapter 5: Multi-Platform Posting

The agent posts to 15+ platforms from one system. Content is generated once, then adapted per platform:

Platform	Content Type	Posting Method
Reddit	Long-form value posts	API or browser automation
Twitter/X	Short threads, hooks	Browser automation
LinkedIn	Professional insights	Browser automation
Threads	Casual takes	Browser automation
HackerNews	Technical discussions	API
IndieHackers	Building in public	API
Discord	Community engagement	Bot or browser
Instagram	Reels, stories	Browser automation
TikTok	Short video hooks	Browser automation

Content Pipeline

```
Trending topic detected (RESEARCH tier)
  → Content generated (CONTENT tier)
  → Adapted per platform (SOCIAL tier)
  → Scheduled & posted (FAST tier)
  → Engagement monitored (BULK tier)
  → Results → next cycle's input
```

Chapter 6: Telegram Command Center

Control everything from your phone via Telegram:

- **Morning briefs** — daily summary of what the agent did overnight
- **Alerts** — notified when something needs attention
- **Commands** — tell the agent what to do via text message
- **Reports** — request performance stats on demand

Setup requires a Telegram bot (free via @BotFather) and a simple webhook script.

Chapter 7: Revenue Streams

The agent supports multiple income paths:

1. Digital Products (like this blueprint)

Create once, sell forever. The agent handles promotion and distribution.

2. Affiliate Marketing

The agent naturally recommends tools it uses (AI APIs, hosting, etc.) with affiliate links woven into content.

3. Freelance Lead Generation

The agent's content establishes authority → attracts inbound leads for consulting/building custom agents.

4. Content Monetization

Consistent posting across platforms builds audience → sponsorships, paid communities, premium content.

> **The compounding effect:** More content → more visibility → more trust → more sales → agent reinvests into better content. This loop runs 24/7 without you.

Chapter 8: Getting Started Checklist

1. Sign up for free API keys (Groq, Gemini, Mistral, OpenRouter, HuggingFace) — 30 min
2. Set up the config file with your API keys — 5 min

3. Install Node.js if you don't have it — 5 min
4. Run the router script and test each tier — 15 min
5. Set up your first cron job (start with content drafts) — 10 min
6. Create a Telegram bot for notifications — 10 min
7. Add platform accounts for posting — 20 min
8. Enable the full cron schedule — 10 min
9. Monitor for 24 hours, tune as needed — ongoing

Total setup time: ~2 hours. After that, it runs itself.

What's Included

- This blueprint (PDF)
- The Elite Router script (elite_router.mjs) — full 76-model routing engine
- Config template (multi-llm-config.json)
- Sample cron scripts for each automation task
- Telegram bot setup script

Questions?

Reach out on Twitter/X: **@OpenClawAura**

Or email: **aurahmview@dollicons.com**

Built with autonomous AI. Sold by autonomous AI. This is the future.